

Ex. 1

```
# Import image
lena = cv2.imread('sift_images/lena.bmp')

# Convert to gray-scale
lena_gray = cv2.cvtColor(lena, cv2.COLOR_BGR2GRAY)

# Define kernels
x_filter = np.array([[-1, 0, 1], [-2, 0, 2], [-1, 0, 1]])
y_filter = np.array([[1, 2, 1], [0, 0, 0], [-1, -2, -1]])

# Compute G_x and G_y
grad_x = cv2.filter2D(lena_gray, cv2.CV_64F, x_filter)
grad_y = cv2.filter2D(lena_gray, cv2.CV_64F, y_filter)

# Compute G_x and G_y (using Sobel operator)
grad_x_sob = cv2.Sobel(lena_gray, cv2.CV_64F, 1, 0)
grad_y_sob = cv2.Sobel(lena_gray, cv2.CV_64F, 0, 1)

# Check differences
diff_x = grad_x - grad_x_sob
diff_y = grad_y - grad_y_sob

print(f"G_x (filter2D) - G_x (Sobel): {diff_x.sum()}")
print(f"G_y (filter2D) - G_y (Sobel): {diff_y.sum()}")

# Absolute magnitude
mag = np.sqrt(grad_x**2 + grad_y**2)

# Convert images to 8-bit
grad_x = cv2.convertScaleAbs(grad_x)
grad_y = cv2.convertScaleAbs(grad_y)
mag = cv2.convertScaleAbs(mag)

plt.figure(figsize=(20,25))
_ = plt.subplot(1,4,1), plt.imshow(lena_gray, cmap=plt.cm.gray), plt.title('Input image')
_ = plt.subplot(1,4,2), plt.imshow(grad_x, cmap=plt.cm.gray), plt.title('$G_x$')
_ = plt.subplot(1,4,3), plt.imshow(grad_y, cmap=plt.cm.gray), plt.title('$G_y$')
_ = plt.subplot(1,4,4), plt.imshow(mag, cmap=plt.cm.gray), plt.title('$G$')
```

Ex. 2.1

```
# Calculate the center of the image
query_image = aachen_day[450:1000, 200:900]
(h, w) = query_image.shape[:2]
(cX, cY) = (w//2, h//2)

# Rotate and scale query image
```

```

M = cv2.getRotationMatrix2D((cX, cY), 45, 2.0)
db_image = cv2.warpAffine(query_image, M, (w, h))
# Display image
fig = plt.figure(figsize=(15, 20))
_ = plt.subplot(1,2,1), plt.imshow(query_image), plt.title('Query image')
_ = plt.subplot(1,2,2), plt.imshow(db_image), plt.title('Database image')
plt.show()

# Detect and compute keypoints and descriptors
query_keypoints, query_descriptors = sift.detectAndCompute(query_image, None)
db_keypoints, db_descriptors = sift.detectAndCompute(db_image, None)
print(f"\nNumber of keypoints (query image): {len(query_keypoints)}.\n")
print(f"Number of keypoints (database image): {len(db_keypoints)}.\n")

# # Distance matrix (rows correspond to query descriptors while columns to database descriptors)
# dist_matrix = np.zeros((len(query_descriptors), len(db_descriptors)))
# for i, desQ in enumerate(query_descriptors):
#     for j, desT in enumerate(db_descriptors):
#         dist_matrix[i, j] = np.linalg.norm(desQ-desT)
# More efficient implementation to compute pairwise distances
dist_matrix = sklearn.metrics.pairwise_distances(query_descriptors, db_descriptors)

# Find the closest database descriptor to each query descriptor
min_distances = dist_matrix.min(axis = 1)
db_descriptors_idx = np.argmin(dist_matrix, axis=1)

# Discard matches above the threshold
thrs = 25
query_descriptors_idx = np.where(min_distances <= thrs)[0]

# Select keypoints, descriptors and indices of retrieved matches
query_keypoints = [query_keypoints[idx] for idx in query_descriptors_idx]
query_descriptors = [query_descriptors[idx] for idx in query_descriptors_idx]
db_descriptors_idx = [db_descriptors_idx[idx] for idx in query_descriptors_idx]
db_keypoints = [db_keypoints[idx] for idx in db_descriptors_idx]
db_descriptors = [db_descriptors[idx] for idx in db_descriptors_idx]
min_distances = [min_distances[idx] for idx in query_descriptors_idx]

print(f"Number of matches: {len(min_distances)}")
# Print Euclidean distance, query descriptor index and database descriptor index of retrieved matches
for k in range(len(min_distances)):
    print(f"Euclidean distance: {min_distances[k]:.2f}, Query Descriptor n. {query_descriptors_idx[k]}, Database Descriptor n. {db_descriptors_idx[k]}")

```

Ex. 2.2 (Optional)

```
row_min = [500, 400, 300, 200, 0]
row_max = [700, 800, 900, 1000, 1063]
col_min = [700, 500, 300, 200, 0]
col_max = [900, 1100, 1200, 1400, 1600]

# List of execution times
exec_time = []

# List of sizes of cropped images
images_size = []

# Threshold matches
thrs = 25

for r_min, r_max, c_min, c_max in zip(row_min, row_max, col_min, col_max):
    start_time = time.time()

    # Calculate the center of the image
    query_image = aachen_day[r_min:r_max, c_min:c_max]
    images_size.append(query_image.shape[:2])
    (h, w) = query_image.shape[:2]
    (cX, cY) = (w//2, h//2)

    # Rotate and scale query image
    M = cv2.getRotationMatrix2D((cX, cY), 45, 2.0)
    db_image = cv2.warpAffine(query_image, M, (w, h))

    # Detect and compute keypoints and descriptors
    query_keypoints, query_descriptors = sift.detectAndCompute(query_image, None)
    db_keypoints, db_descriptors = sift.detectAndCompute(db_image, None)
    print(f"\nNumber of keypoints (query image): {len(query_keypoints)}.\n")
    print(f"Number of keypoints (database image): {len(db_keypoints)}.\n")

    # Pairwise distances
    dist_matrix = sklearn.metrics.pairwise_distances(query_descriptors, db_descriptors)

    # Find the closest database descriptor to each query descriptor
    min_distances = dist_matrix.min(axis = 1)
    db_descriptors_idx = np.argmin(dist_matrix, axis=1)
```

```

# Discard matches above the threshold
query_descriptors_idx = np.where(min_distances <= thrs)[0]

# Select keypoints, descriptors and indices of retrieved matches
query_keypoints = [query_keypoints[idx] for idx in query_descriptors_idx]
query_descriptors = [query_descriptors[idx] for idx in query_descriptors_idx]
db_descriptors_idx = [db_descriptors_idx[idx] for idx in query_descriptors_idx]
db_keypoints = [db_keypoints[idx] for idx in db_descriptors_idx]
db_descriptors = [db_descriptors[idx] for idx in db_descriptors_idx]
min_distances = [min_distances[idx] for idx in query_descriptors_idx]

# add exection time
exec_time.append(time.time() - start_time)

_ = plt.figure(figsize=(6,6))
plt.plot(range(1, len(row_min)+1), exec_time, '--o'), plt.title('Execution time'), plt.grid('on')
# Integer x-axis
plt.locator_params(axis='x', integer=True, tight=True)
plt.xlabel('Experiment'), plt.ylabel('Time [s]')
for i in range(len(row_min)):
    plt.text(i, exec_time[i], f" ({images_size[i][0]}, {images_size[i][1]})")

```

Ex. 3

```

# Calculate the center of the image
query_image = aachen_day[450:1000, 200:900]
(h, w) = query_image.shape[:2]
(cX, cY) = (w//2, h//2)
# Rotate and scale query image
M = cv2.getRotationMatrix2D((cX, cY), 45, 2.0)
db_image = cv2.warpAffine(query_image, M, (w, h))

# Detect and compute keypoints and descriptors
query_keypoints, query_descriptors = sift.detectAndCompute(query_image, None)
db_keypoints, db_descriptors = sift.detectAndCompute(db_image, None)

# Pairwise distances
dist_matrix = sklearn.metrics.pairwise_distances(query_descriptors, db_descriptors)

## K-NN matcher (K=2)
# For each query descriptor find the indices and distances

```

```

# of the 2 closest database descriptors.

k = 2
distances = np.partition(dist_matrix, kth=k, axis=1)[:, :k]
db_idx = np.argpartition(dist_matrix, kth=k, axis=1)[:, :k]

# Apply ratio test
min_dist = []
min_query_idx = []
min_db_idx = []
for k in range(distances.shape[0]):
    if distances[k][0] < 0.75 * distances[k][1]:
        min_dist.append(distances[k][0])
        min_query_idx.append(k)
        min_db_idx.append(db_idx[k][0])

print(f"Number of query keypoints: {len(query_keypoints)}")
print(f"Number of database keypoints: {len(db_keypoints)}")
print(f"Number of matches: {len(min_dist)}")

# Print first 10 matches
nMatches = 10
for k in range(nMatches):
    print(f"Euclidean distance: {min_dist[k]:.2f}, Query Descriptor: n. {min_query_idx[k]}, Database Descriptor: n.{min_db_idx[k]}")

```

Ex. 5

```

# Brute-force matcher
bf = cv2.BFMatcher_create(cv2.NORM_L2)

# Finds the best match for each descriptor from a query set.
matches = bf.match(query_descriptors, db_descriptors)

# Order matches
ordered_matches = sorted(matches, key = lambda m:m.distance)

# Discard matches above the threshold
thrs = 100

matched_des = [ordered_matches[idx] for idx in range(len(ordered_matches)) \
               if ordered_matches[idx].distance <= thrs]

```

```
print(f"Number of matches: {len(matched_des)}")
print(f"Min distance: {matched_des[0].distance}, Max distance: {matched_des[-1].distance}.")

fig = plt.figure(figsize=(15,10))
draw_params = dict(matchColor=(0, 255, 0), singlePointColor=(255, 0, 0), flags=0)
out_matches = cv2.drawMatches(query_image, query keypoints, database_image, db keypoints,
matched_des[:50], None, **draw_params)
_ = plt.imshow(out_matches)
```